

Ryan Dielhenn

13112 Valleyheart Dr., APT 303, Studio City, CA 91604

📞 (818) 519-6414 | 📩 dielhennr@gmail.com | 🌐 ryandielhenn.github.io

Technical Skills

| | |
|-----------------------|--|
| Languages | Go, Rust, Java, C, Python, Scala, SQL |
| Infrastructure | Docker, Linux Systems Administration, Prometheus, Grafana, Git, Kafka, etcd, DuckDB, Scripting/Automation (Python, Bash) |
| Networking | TCP/IP, UDP, BGP, OSPF, DNS, DHCP, Gossip Protocols, Phi-accrual Failure Detection |
| Concepts | Distributed Systems, High-Performance Data Processing, Network Observability, Cluster Replication, API Design |

Experience

Confluent

Mountain View, CA

SOFTWARE ENGINEER

- Contributed to Apache Kafka's migration from ZooKeeper to KRaft, improving observability and reliability of the distributed consensus layer in Java/Scala.
- Built metrics pipelines in Java/Scala to monitor cluster health, quorum state, and inter-broker communication patterns for KRaft; updated Confluent Cloud tooling in Go to integrate new metrics.
- Contributed to Confluent Cloud's Cluster Linking integration with KRaft architecture for cross-cluster replication.

Confluent

Mountain View, CA

SOFTWARE ENGINEERING INTERN

- Implemented dynamic client reconfiguration in Java/Scala for Apache Kafka, enabling runtime updates to producer/consumer settings (including connection, security, retry, and ack configurations) without service restarts.
- Enhanced Confluent Cloud's rebalance tooling with asynchronous replica movement support.
- Continued contributing to Apache Kafka during Fall 2020 while completing undergraduate degree.

University of San Francisco

San Francisco, CA

UNDERGRADUATE RESEARCHER

- Developed a C implementation of Geopresence, a bitmap-based geospatial indexing system using RoaringBitmap compression and HyperLogLog++ for efficient location queries on IoT devices.
- Implemented C-based adaptive grid index achieving 17x speedup over Java; outperformed R-tree indexing by 400x at scale (1M+ points) while R-trees showed advantages on sparse datasets (<7K points).

Projects

SlateDB & OpenData

OPEN SOURCE CONTRIBUTOR

February 2026 – Present

- Contributing to SlateDB, an embedded key-value storage engine on object storage; simplified the builder API by consolidating compactor configuration, improving developer experience.
- Contributing to OpenData, a collection of databases sharing a unified object storage foundation via SlateDB; implemented TSDB benchmarks, added scalar argument support for PromQL expressions, and Prometheus metadata API support.

ZephyrCache

DISTRIBUTED SYSTEMS PROJECT

July 2025 – Present

- Built a distributed cache in Go with consistent hashing for request routing and etcd for dynamic node registration, service discovery, and lease-based failure detection.
- Achieved 49k ops/sec with 32 concurrent clients over HTTP, benchmarked on a virtual Docker bridge network with inter-node latency monitoring.
- Deployed a Prometheus/Grafana observability stack to monitor network throughput, HTTP latency distributions, and container resource utilization across nodes.

Education

California State University, Los Angeles

Los Angeles, CA

M.S. IN COMPUTER SCIENCE (IN PROGRESS), GPA: 4.0

Expected 2026

- Thesis:** Machine Learning Approaches to Trust and Reputation in Networked Environments.
- Developing a bare-metal Raspberry Pi cluster to evaluate Byzantine Fault Tolerance using UDP-based Gossip protocols and Phi-accrual failure detection.
- Relevant Coursework: Advanced Computer Networking (BGP, OSPF, TCP/IP), Software Engineering, Artificial Intelligence, Machine Learning.

University of San Francisco

San Francisco, CA

B.S. IN COMPUTER SCIENCE, MINOR IN MATHEMATICS, GPA: 3.75

2016 – 2020

- Relevant CS Coursework: Operating Systems, Computer Architecture, Data Structures & Algorithms, Big Data, Software Development, Programming Language Paradigms, Senior Capstone
- Relevant Math Coursework: Linear Algebra, Formal Methods (Logic and Proof), Calculus I & II, Modern Algebra