

Ryan Dielhenn

13112 Valleyheart Dr., APT 303, Studio City, CA 91604

📞 (818) 519-6414 | 📩 dielhennr@gmail.com | 🌐 ryandielhenn.github.io

Education

California State University, Los Angeles

M.S. IN COMPUTER SCIENCE (IN PROGRESS)

- Relevant Coursework: Advanced Artificial Intelligence, Advanced Software Engineering, Data Science

Los Angeles, CA

Expected 2026

University of San Francisco

B.S. IN COMPUTER SCIENCE, MINOR IN MATHEMATICS, GPA: 3.75

- Relevant CS Coursework: Big Data, Software Development, Data Structures & Algorithms, Operating Systems, Computer Architecture, Programming Language Paradigms, Senior Capstone
- Relevant Math Coursework: Calculus I & II, Formal Methods, Linear Algebra, Abstract Algebra

San Francisco, CA

2016 – 2020

Projects

ZephyrCache

DISTRIBUTED CACHING SYSTEM

- Designing and building a distributed caching system in Go using consistent hashing for key distribution across nodes.
- Integrating etcd for cluster membership, lease management, and peer discovery to enable node coordination.
- Developing monitoring and benchmarking tools to measure routing efficiency and system performance.

Los Angeles, CA

August 2025 - Present

EDA Dashboard

HIGH PERFORMANCE DATA ANALYSIS TOOL

- Built a web-based EDA tool with a FastAPI backend and Streamlit frontend for visualizing and analyzing large datasets.
- Integrated DuckDB for fast, in-process SQL queries enabling responsive exploration of large datasets.
- Implemented bias detection and data drift monitoring features for ML pipeline observability.

Los Angeles, CA

September 2025 - Present

Experience

Confluent

SOFTWARE ENGINEER

- Collaborated across engineering teams to improve reliability, observability, and security during Kafka's transition to a ZooKeeper-free architecture (KRaft).
- Adapted Cluster Linking to support KRaft, enabling cross-cluster replication without ZooKeeper.
- Integrated metrics pipelines to monitor cluster health and quorum state in KRaft mode.

Mountain View, CA

January 2021 – July 2022

Confluent

SOFTWARE ENGINEERING INTERN

- Improved Kafka usability as a Summer 2020 intern by implementing dynamic client reconfiguration, allowing clients configurations to be updated without restarting producer and consumer clients.
- Enhanced Confluent Cloud's rebalance tooling with asynchronous replica support.
- Continued contributing to Apache Kafka during Fall 2020 while completing my undergraduate degree.

Mountain View, CA

May 2020 – August 2020

University of San Francisco

RESEARCH ASSISTANT

- Conducted research on distributed systems and edge computing architectures under faculty supervision.
- Designed and implemented a geospatial indexing system (Geopresence) optimized for IoT and low-power devices, using RoaringBitmap for bitmap compression and HyperLogLog++ for approximate cardinality estimation.
- Benchmarked prototypes and demonstrated hyper-local, location-aware queries (e.g., retrieving air quality data directly from nearby sensors instead of centralized APIs).

San Francisco, CA

January 2020 – December 2020

TEACHING ASSISTANT — BIG DATA & OPERATING SYSTEMS

- Led weekly office hours, provided project design/debugging support, and evaluated student assignments.

August 2019 – May 2020

ASSISTANT SYSTEMS ADMINISTRATOR

- Automated updates and maintenance tasks for Linux lab machines, reducing manual overhead for IT staff.
- Diagnosed and resolved hardware/software issues for faculty and students in a high-demand academic environment.

May 2019 – August 2019

Technical Skills

Languages Go, Python, Java, Scala, SQL

Frameworks FastAPI, Streamlit

Infrastructure Kafka, Docker, etcd, Git, DuckDB

Concepts Distributed Systems, API Design, Data Modeling, Performance Optimization